

Data Mining: Final Project

Graded checkpoint due: **Tuesday 4/17** before 5:00 PM via Moodle

Final code and paper due: **Thursday 5/10** before 5:00 PM via Moodle

1 Introduction

For this project, we will participate in the Kaggle.com competition “House Prices: Advanced Regression Techniques”. Kaggle.com is a website that hosts data science competitions – some for significant cash prizes. It’s pronounced “Kaa-gle”, with the first syllable rhyming with “wag” like a dog’s tail. (When I first learned of the website, people around me were calling it “Kay-gle”, so I may mess this up as I try to change my pronunciation.)

Participating in Kaggle competitions is a nice thing to put on your resume; I’ve seen some job advertisements mention such competitions explicitly, in fact. Go to Kaggle.com and create an account if you don’t have one already.

Log in, and then go here for the contest we’ll participate in:

<https://www.kaggle.com/c/house-prices-advanced-regression-techniques>

Skim around on those pages a bit to get a sense for the competition and the problem domain. Under the data tab, there are various files you can download. There’s no need to download the three .gz files – those are just zips of some of the other files, which are small enough anyway. Grab all the other files, though.

1.1 Reading the Rest of this Document

I know there’s a lot of text here, but just like in the assignments, it’s here for your benefit if you take the time to consider every detail. You can learn a lot just from studying this document, so I want to encourage you to treat it like a textbook reading assignment – underline, take notes, summarize key ideas in your own words, etc.

1.2 What’s the Scope of this Project?

Briefly, in this project you will process this dataset, apply some machine learning algorithms to it, tune your processing and algorithms through many experiments, and write a formal report of your work. This is a significant project, accounting for a significant part of your final grade in the course: somewhere around 50 points, around the same weight as the midterm and final exams.

In class we will continue to cover additional topics throughout the semester. Therefore, some of your out-of-class time will be spent studying in-class material. Of course there will be many variations among people in the class and from week to week. But I might expect that on average for the rest of the semester, you’d spend about 2 hours per week working on the ongoing class material, and another 5 hours per week working on the project. This makes just a bit over 2 hours out of class for every hour in, which is a fair target.

So to summarize, the project should represent about 5 hours per week of work. Times 5 weeks, that makes about 25 hours of work that should be reflected in the project. This is 25 hours, assuming a fair amount of comfort with the course material. If you are behind in your understanding of course material up to this point, more time may be required, to catch up.

If you're interested in taking things even further, though, you might consider working even more, as much as you can and would like. This would be a great item to add to your resume. This work may also inspire you for a very interesting senior project, if you haven't yet done the senior project.

1.3 Using Code and Ideas You Find Online

There are a lot of good ideas, and a lot of code, online. You are welcome to read up on whatever you want and use whatever you want. Done correctly, this can be a great learning experience. If you copy code (whether a single line or many lines) from somewhere, you must provide BEGIN, END, EXPLANATION, and ADDED comments as illustrated below:

```
def printS(s):
    # BEGIN: from http://www.stackoverflow.com/fake
    # EXPLANATION: Loops through s one character at a time.
    #   -1 is the index for the last character,
    #   so s[:-1] means to slice up to but not including last char
    for c in s[:-1]:
        print(c, end='-')
    print(s[-1]) # ADDED: To handle last character differently
    # END: from http://www.stackoverflow.com/fake
```

If you can't explain the code, I'm asking you not to use it. If you make adjustments to the code, indicate that with "ADDED" comments. If you didn't copy any code, but just got an idea from somewhere, just cite the website.

You must have at least 50% of the code be your own. That means that if you copy a lot, then hey, that's cool. You just need to write more code of your own to build off of that foundation you got from online.

I know what material is out there, and it's easy for me to quickly google things in your code we haven't talked about in class. **If you use material you find online but do not cite it as described above, there will be severe penalties, up to and including failing the course and being reported for academic dishonesty.** If you have any questions or concerns, please feel free to ask me *before you turn in your work* – no problem, we'll work it out together! This is a major issue in real-life work. For example, do you know about the lawsuit between Oracle and Google over copyright infringement of comparatively very small snippets of code? ([https://en.wikipedia.org/wiki/Oracle America, Inc. v. Google, Inc](https://en.wikipedia.org/wiki/Oracle_America, Inc. v. Google, Inc)) It's still going on! (<https://www.wired.com/story/the-case-that-never-ends-oracle-wins-latest-round-vs-google/>)

1.4 Working with Other People in Class

You should do everything with your partner, of course. Other than that, please don't share any code with anyone, or look at anyone else's code from this course or previous offerings of the course. You are welcome to discuss big-picture ideas without code, though. As with the online resources discussion above, if you have questions or concerns, please talk with me *before* doing it.

2 Completing this Project with Excellence

2.1 Data Mining: An Iterative Process with Feedback Loops

As described in detail below, your work will include processing the data, applying algorithms, experimenting and tuning your approaches, and writing a formal report of your results. However, do not think of these as sequential steps. Rather, this is an iterative process, where lessons learned in a “later” step will lead to more enhancements in “earlier” steps.

For example, you might do some initial pre-processing of the data and apply some algorithms, only to get mediocre results. You write up these results, and in your write-up, hypothesize about why the results were not very good. The act of writing takes time, but it forces you to think in more detail about what you’ve done, clarifying your thinking and leading to new ideas. You implement these new ideas, experiment some more, write up some more, etc.

In the following sections, I provide more advice about the kinds of things an excellent (A-level) project will include.

2.2 Pre-process the Data

This step is called “pre-processing” because it’s necessary before we attempt to build a model. This is not to suggest that this work happens only once, at the beginning of the project, though. As described above, you will need to refine your pre-processing throughout the project as you obtain experimental results and try additional ideas.

The first step is to study your dataset. Know it deeply! Pore over it! If you don’t know all about what your data says, you can’t use it effectively. Start with `data_description.txt`. It describes the possible values for each attribute. For each attribute, ask yourself:

- What is the type of this attribute? (Nominal, ordinal, interval, ratio)
- Do I need to convert its type to something else? If so, what?
- Do I need to normalize or standardize this attribute?
- Are there missing attributes? How many? How should I handle them? Should I replace missing values? Delete rows with missing values? Delete columns with missing values? How is a missing value indicated for this attribute? Is a value truly “missing”, or can I derive it somehow from other values? Does missing mean “unknown”, or “not applicable”, or what, and how might this affect how I handle that “missing” value?
- Intuitively, do I expect that this attribute is very important, somewhat important, or not very important? In terms of correlation with other attributes (using the `corr` method), is this attribute highly correlated with others, meaning that some should be dropped?
- Is the attribute highly skewed, therefore possibly suggesting that a log transformation would be useful? (You might want to Google this.)
- Should I transform the values in some way to more clearly capture what is important about it?
- What do I understand intuitively about this attribute that should influence how my system uses it? What do I expect to be the case about this attribute, and how can I explore whether or not that expectation is correct?

Note that your answers to the above questions may be different for each attribute. Several attributes may behave similarly, but if you think they all behave the same way, you probably have not considered them deeply enough.

You might also ponder whether it'd be useful to create additional attributes. Might a new attribute be derived from existing attributes somehow? What key ideas could be determined from the data that are not explicitly represented from the data at this time?

Plan on spending lots of time on pre-processing – quite possibly more time here than on any other step. It makes an enormous difference in the success of your model! **A significant part of your project's success (and your final grade) will depend on how carefully you consider questions like those above, and how well you write up your actions and attempts.**

As you make pre-processing decisions, be certain you keep a careful log of what you're doing. Keep writing things up as you go, to help you organize your thoughts, and so you don't forget what you've done and why later on. You may think you'll remember all your decisions and explain them well later, but... you won't! There's just too much to do to keep it all in your head week after week.

2.3 Building a Model

With some initial pre-processing done, you're ready to feed your data into an algorithm to create a model. Using the built-in classes of scikit-learn, it's surprisingly easy to try an algorithm out. Anything about regression is a potentially applicable algorithm to this task, since this is a regression task (as opposed to a classification task). The Kaggle competition description, for example, mentions specifically "gradient boosting". So please start by looking into the GradientBoostingRegressor class.

It may take you a little time to figure out how to create and use a built-in model the first time. Note, however, that you've already done this in hw08 with KNeighborsClassifier. While this is a classifier rather than a regressor, for the most part they are used in the same way – just on different types of data.

If your code crashes when you try to apply an algorithm, it may be that some kind of pre-processing is necessary that you haven't done. For example, does your algorithm require all numerical attributes?

After you get some kind of at least not terrible performance with GradientBoostingRegressor, try several other algorithms. I'd encourage you to do a Google search on things like "scikit-learn regressor". This should be comparatively quick and easy to do given what you've already figured out. You basically just call a constructor, and pass the object to cross_val_score – the same as with KNeighborsClassifier in hw08, and the same as with GradientBoostingRegressor.

Unfortunately, we will not have time to learn how all of these different regression algorithms work. If you were doing this as a career, you would need to spend more time understanding the algorithms and determining which best apply to your problem. We can still make progress, however, by treating them as black boxes.

As you do this, make sure you're spending time organizing and cleaning up your code! Good design, good variable and function names, avoiding code repetition, etc. – these efforts are all time well-spent, since you're in this for the long haul.

As you begin to get results from these algorithms, review your pre-processing work again and determine if there's something you want to adjust at this time.

2.4 Tune Your Algorithms and Pre-Processing Via Experiments

To “tune” your work means to:

1. Hypothesize what adjustments may be useful:
 - a. Should you pre-process some data differently?
 - b. As you read the documentation about a particular algorithm in more detail, what parameters does the algorithm have? Which ones seem most important? Might it make sense to use a value other than the default for some of them?
 - c. Is there a different algorithm altogether that you should try?
2. Make the adjustments and run experiments on the resulting system performance.
3. Log your results, and interpret them. Try to explain them. Return to step (1).

One thing you may consider trying is the advice here, as an example of tuning gradient boosting:

<https://www.analyticsvidhya.com/blog/2016/02/complete-guide-parameter-tuning-gradient-boosting-gbm-python/>

You can follow similar principles for other algorithms, and also see what else you find online.

It is very important that you document carefully, in your code, in Excel spreadsheets, and in text notes, what experiments you're trying and what results you get. Again, you may think you'll remember what's what, but very quickly things will become disorganized and confusing. You'll forget what results correspond to what system configurations, your experiments will be inadequately controlled, etc. I speak from experience, both in my own past (!), and in work with other students.

2.5 Write a Formal Report of Your Results

As I've advised above, please be certain that you're writing throughout your project, not just at the end. The writing will help you clarify your thoughts and determine where to go next as you work. It will also help you get important ideas down while they're fresh in your mind, before you forget some details.

To give you some experience in another important technical skill, every paper must be written in LaTeX. Please see the separate pdf on LaTeX for details.

I will provide precise content requirements for the report very soon.

3 Significant Computing Power to Help You!

I'm excited to announce that I have received a grant for over \$15,000 worth of access to advanced cloud computing infrastructure funded by the NSF, for our use in this project. This is awesome because:

- It will enable you to run the experiments you want to run, with fewer limitations based on computational cost.
- It will give you experience in managing parallel execution of your code, working with cloud resources, and working in a Unix environment.

It's important that we share this access and manage it carefully, or we will run out and wreck things for everyone. I will provide training for this access soon.

4 Deadlines

There are two deadlines for this project:

1. A graded checkpoint submission due about halfway through.
2. The final products due at the end of the semester.

Please see the top of this document for precise dates and times of these deadlines.

4.1 The Graded Checkpoint

For the graded checkpoint, you must have the following completed:

- a) Basic data pre-processing – enough to be able to run an algorithm while using at least 50% of the attributes.
- b) At least one model created and results obtained.
- c) A Word document clearly describing your pre-processing steps, the algorithm(s) you're running, and the results you've obtained so far.

Turn in a zip of your code and your Word document. Name the zip *lastname1-lastname2-checkpoint.zip*.

4.2 The Final Products

For the final product, please have the following completed:

- a) More refined data pre-processing.
- b) Multiple models created: different algorithms, different pre-processing strategies, and different parameterizations, all resulting from your tuning processes.
- c) A LaTeX document clearly describing all of this work, including your pre-processing, hypotheses, experiments, and results. Figures, graphs, and tables should be used when appropriate for capturing your work.

Turn in a zip of your code and all LaTeX documents (.tex, .pdf, figures, generated files like .out, everything). Name the zip *lastname1-lastname2-final.zip*.

5 Take Advantage of this Opportunity

Beyond motivation of grades, graduating, etc., I want to encourage you, as much as you're able and interested, to go all out on this project. This could be a wonderful showcase of your abilities as you apply for jobs, and can stretch your critical thinking abilities perhaps farther than ever before – ***if*** you decide to take advantage of this opportunity. I am eager to work with you in this process.